

More Agents Isn't Always Better: across 4 NLP domains and 2 base models, multi-agent scaling produces three different curve shapes. A simple threshold on knowledge concentration K predicts the optimal agent count for 7/8 trustable cells (88%, leave-one-domain-out CV).

